

What did you learn from Vtune profiling:

I've used Vtune Threading option which is a feature that allows you to analyze the performance of a multi-threaded application. When analyzing a multi-threaded application, VTune Threading enables you to gather performance data from each thread separately and then aggregate the results to get a more comprehensive view of the application's performance.

With this option, I've found places with hotspots, which are functions on the project, that were using the most time in the application and the exact duration spent in each function.

This information gave me the performance bottleneck in the code and where should I focus on.

From the profiler I've known that the best place to perform optimization was `calculate_derivatives`, and a couple more functions which I'll mention later in the report.

Each one of these hotspots that were found, I dealt with them with the `openmp` pragmas of `fortran`.

Picture from the Vtune showing the functions and the time each one spent on the run: (We can see that `calculate derivatives` took the most time, so we started by handling it first)

Grouping: Module / Function / Call Stack				
Module / Function / Call Stack	CPU Time			Idle
	Effective Time ▼	Spin Time	Overhead Time	
▼ main	41.875s	0s	0s	
▶ __velocity_module_MOD_calculate_derivatives	19.799s	0s	0s	
▶ __vertex_mass_module_MOD_calculate_vertex_mass	5.493s	0s	0s	
▶ __hydro_step_module_MOD_calculate_acceleration_3d	3.703s	0s	0s	
▶ __hydro_step_module_MOD_calculate_thermodynamic	2.456s	0s	0s	
▶ __volume_module_MOD_calculate	2.364s	0s	0s	
▶ __hydro_step_module_MOD_calculate_artificial_viscosi	1.452s	0s	0s	
▶ __hydro_step_module_MOD_calculate_velocity_3d	1.216s	0s	0s	
▶ __hydro_step_module_MOD_calculate_energy_3d	1.016s	0s	0s	
▶ __ideal_gas_module_MOD_calculate_ideal_gas	0.974s	0s	0s	
▶ __rezone_module_MOD_calculate_rezone_3d	0.780s	0s	0s	
▶ __hydro_step_module_MOD_calculate_thermodynamic	0.452s	0s	0s	
▶ __time_module_MOD_calculate_vel_grad_dt_3d	0.444s	0s	0s	
▶ __hydro_step_module_MOD_calculate_density_avg_3d	0.352s	0s	0s	

(I have to mention that the Vtune profiler didn't work as I expected since I have the problem of seeing the function as there location on the code, assembly location, I informed Yehonathan about the issue.)

My achievements:

First thing I did was running the project on serial mode without any optimization, to check the amount of time the running would take, and I've got this:

(Sorry for the print("hi") that were in the middle)

```
Application Output
0
  making sedov-taylor
done mesh
Done diagnostics
finished building problem
hi
Cycle time:    12.693916454911232
hi
Cycle time:    13.024112544953823
hi
Cycle time:    12.779446320608258
hi
Cycle time:    12.647634431719780
hi
Cycle time:    12.583921505138278
hi
Cycle time:    12.511623922735453
hi
Cycle time:    12.472357081249356
hi
Cycle time:    12.476540770381689
hi
Cycle time:    12.430432088673115
hi
Cycle time:    12.474860582500696
Total Time:    126.13239316642284
ncyc:          10
```

After that, I've decided to use -o3 flag on the cmake list, which is a compiler optimization flag that enables the highest level of optimization in the compiler.

And I've got the result of **50 seconds** total time:

```
u187554@s001-n002:~/helperScal/ScalSALE-main/src/Scripts$ ./run.sh 1
0
  making sedov-taylor
done mesh
Done diagnostics
finished building problem
Cycle time:    5.1485561244189739
Cycle time:    5.0343927331268787
Cycle time:    5.0332871153950691
Cycle time:    5.0297751836478710
Cycle time:    5.0369162596762180
Cycle time:    5.0265680160373449
Cycle time:    5.0231137685477734
Cycle time:    5.0301657766103745
Cycle time:    5.0321176201105118
Cycle time:    5.0258645042777061
Total Time:    50.449725050479174
ncyc:          10
u187554@s001-n002:~/helperScal/ScalSALE-main/src/Scripts$
```

Last thing, was doing the optimizations using the open-mp pragmas of fortran while keeping -o3 flag, and I've reached the total time of 10 second approx.

```
u187554@s001-n061:~/ScalSALE-main/src/Scripts$ ./run.sh 1
0
making sedov-taylor
done mesh
Done diagnostics
finished building problem
Cycle time: 1.3682540468871593
Cycle time: 1.0553217660635710
Cycle time: 1.0775032900273800
Cycle time: 1.1221211627125740
Cycle time: 0.94899427518248558
Cycle time: 1.0067895520478487
Cycle time: 0.91284256055951118
Cycle time: 0.91957117430865765
Cycle time: 0.89623875729739666
Cycle time: 0.94888567738234997
Total Time: 10.332222618162632
ncyc: 10
u187554@s001-n061:~/ScalSALE-main/src/Scripts$
```

The places where I used OpenMp pragmas:

To parallelize a specific part of my application using OpenMP, I first identified the section of code suitable for parallelization, which happened to be calculate_derivatives in velocity.f90, and so on, based on the previous identification using Vtune.

After parallelizing first function (in velocity.f90), I recompiled the code, and ran Vtune again to check the accomplishments so far, and checked again what's the next place to parallelize. I kept parallelizing function and compiling and checking again, until I saw no further optimization comes out of parallelizing more function, and there I stopped and that's were the results:

We can see in the picture the places where the pragma "omp parallel do" was added, in each place of them there was a function that the vtune gave us, indicating that it ran a good amount of time and could be parallelized.

```
./Material/Equation_of_state/ideal_gas.f90:60:      !$omp parallel do private(k,j,i)
./Time_step/hydro_step.f90:527:      !$omp parallel do collapse(3) private(k,j,i,tmp_mat)
./Time_step/hydro_step.f90:561:      !$omp parallel do private(k,j,i,tmp_mat)
./Time_step/hydro_step.f90:583:      !$omp parallel do private(k,j,i)
./Time_step/hydro_step.f90:600:      !$omp parallel do private(k,j,i)
./Time_step/hydro_step.f90:675:      !$omp parallel do private(k,j,i)
./Time_step/hydro_step.f90:684:      !$omp parallel do private(k,j,i)
./Time_step/hydro_step.f90:821:      !$omp parallel do &
./Time_step/hydro_step.f90:1177:      !$omp parallel do &
./Time_step/hydro_step.f90:1826:      !$omp parallel do &
./Time_step/.ipynb_checkpoints/hydro_step-checkpoint.f90:563:      !$omp parallel do simd collapse(1) schedule(simd:static) private(j,i,tmp_mat)
./Time_step/.ipynb_checkpoints/hydro_step-checkpoint.f90:585:      !$omp parallel do simd collapse(1) schedule(simd:static) private(j,i)
./Time_step/.ipynb_checkpoints/hydro_step-checkpoint.f90:601:      !$omp parallel do simd collapse(1) schedule(simd:static) private(j,i)
./Time_step/.ipynb_checkpoints/hydro_step-checkpoint.f90:680:      !$omp parallel do simd collapse(1) schedule(simd:static) private(j,i)
./Time_step/.ipynb_checkpoints/hydro_step-checkpoint.f90:824:      !$omp parallel do simd collapse(1) schedule(simd:static)
./Time_step/time.f90:323:      !$omp parallel do private(k,j,i,vel_diff_coor_diff_i,vel_diff_coor_diff_j,vel_diff_coor_diff_k,vel_grad)
./Time_step/time.f90:534:      !$omp parallel do private(k,j,i,vel_diff_x,vel_diff_y,vel_diff_z,len_sq, &
./Rezone_and_Advect/rezone.f90:252:      !$omp parallel do private(k,j,i)
./Rezone_and_Advect/rezone.f90:270:      !$omp parallel do private(k,j,i)
./Rezone_and_Advect/rezone.f90:292:      !$omp parallel do private(k,j,i)
./Rezone_and_Advect/rezone.f90:304:      !$omp parallel do private(k,j,i)
./Quantities/Cell/volume.f90:122:      !$omp parallel do collapse(3) &
./Quantities/Vertex/vertex_mass.f90:191:      !$omp parallel do private(k,j,i,t,kk,jj,ii,i1,i2,i3,j1,j2,j3,k1,k2,k3)
./Quantities/Vertex/velocity.f90:233:      !$omp parallel do &
```

Describe what compilers you chose to use:

At first I searched google for what is better between the three compilers: ifort, gfortran and ifx, and I got this:

ifort is known for its high performance and optimization capabilities, particularly when running on Intel processors. This can make it a good choice for applications that require maximum performance, such as scientific or engineering simulations.

And ifort has good support for modern Fortran standards, which can make it easier to write and maintain high-quality Fortran code. It also has good support for OpenMP and MPI, which can be useful for parallelizing code on multi-core or distributed systems.

And after that I tested the three compilers on the second phase of parallelization (Where I had -O3 flag and no further openMP optimization), and I found that the three of them gave the same time, so given the explanation that I found on google, I've decided to go with iFort.

GPU offloading:

1. The program consists of a main loop (in problem.f90) which is dependent on previous iterations, making it difficult to offload to the GPU, To offload a function to the GPU, the longest function was chosen which was Calculate_derivatives (which was found in the vtune before).
2. Offloading the entire program to the GPU would be not feasible for this exercise, so I didn't do that, I only offloaded the calculate_derivative.
3. The solution is to map the relevant data structures before the main loop, update them on each iteration, and transfer arrays from CPU to GPU before Calculate_derivatives, and arrays from GPU to CPU after.

The implementation of this offloading was in two files, the problem.f90 and in velocity.f90, and all this was in the GPU folder.

I've added the relevant pragma and offloading function before the main loop in problem.f90, and add before and after the loop in velocity.f90 (calculate_derivatives) to send the data and get it back and update...